

EECE 5644: Machine Learning Project 2 Report

1. Given the following data:

$$p(\mathbf{x}) = p(\mathbf{x}|L=0)p(L=0) + p(\mathbf{x}|L=1)p(L=1)$$

$$p(L=0) = 0.65$$

$$p(L=1) = 0.35$$

$$p(\mathbf{x}|L=0) = w_1 g(\mathbf{x}|m_{01}, C_{01}) + w_2 g(\mathbf{x}|m_{02}, C_{02})$$

$$p(\mathbf{x}|L=1) = g(\mathbf{x}|m_1, C_1)$$

$$w_1 = w_2 = 0.5$$

class conditional parameters:

$$m_{01} = \begin{bmatrix} 3 \\ 0 \end{bmatrix} \quad C_{01} = \begin{bmatrix} 2 & 0 \\ 0 & 1 \end{bmatrix} \quad m_{02} = \begin{bmatrix} 0 \\ 3 \end{bmatrix} \quad C_{02} = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix} \quad m_1 = \begin{bmatrix} 2 \\ 2 \end{bmatrix} \quad C_1 = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$




Figure 1: True data distribution for the 10k validation dataset

A ERM classification (Bayes Classifier)

- i Minimum expected risk classification rule
Likelihood ratio test:

$$\frac{p(x|L=1)}{p(x|L=0)} \underset{L=0}{\overset{L=1}{>}} \frac{p(L=0)}{p(L=1)} * \frac{(\lambda_{10} - \lambda_{00})}{(\lambda_{01} - \lambda_{11})} = \gamma$$

For 0 – 1 loss classification, the above reduces to:

$$\frac{p(x|L=1)}{p(x|L=0)} \underset{L=0}{\overset{L=1}{>}} = \frac{0.7}{0.3} * \frac{(1-0)}{(1-0)} = \frac{0.65}{0.35}$$

In the above formula, $\gamma = 2.33$ is the theoretical threshold. Classify $L = 1$ if

$$\frac{p(x|L=1)}{p(x|L=0)} > 1.86$$

and classify $L = 0$ otherwise.

- ii ROC curve

The minimum Probability error classifier is implemented on the 10K validation set samples. The discriminant scores were compared to the threshold values (which calculated as the mid-point value of the discriminant scores), and the ROC curve plotted based on this threshold value. From the figure, the black dot indicates the point probability of error is minimum. The true positive and false positive rates at this minimum probability error is 0.55 and 0.15 respectively.




Figure 2: ROC curve with minimum probability error

types	γ	minimum probability error
Theoretical	1.86	0.41
Experimental	2.341	0.25

The following plot shows the distribution after classification:




Figure 3: Data distributions after classification

B MLE for logistic-linear-function-based approximation: This approximator is based on 3 separate training datasets and later implemented on 10K test datasets. The theta parameters were trained by minimizing the cost function using gradient descent optimization technique. The cost function is the negative-log-likelihood of each training data:

$$cost = -\frac{1}{N} \sum_{i=1}^N [y_i \ln(h(x_i, w)) + (1 - y_i) \ln(1 - h(x_i, w))]$$

where, the sigmoid function $h(x_i, w)$ is

$$h(x_i, w) = \frac{1}{1 + e^{-w^T z(x)}}$$

$$z(x) = [1, x^T]^T$$

i Using the 20 training samples:




Figure 4: Classification of 10,000 validation samples based on 20 training samples

ii using the 200 training samples




Figure 5: Classification of 10,000 validation samples based on 200 training samples

iii using the 2000 training samples




Figure 6: Classification of 10,000 validation samples based on 2000 training samples

Training samples	minimum probability
20	0.367
200	0.346
2000	0.343

Observation:

As the number of the training samples increases, the minimum probability error decreases. Also as the number of training samples increases, we can see from the above figures that, the mis-classification labels decreases (which is shown in red labels).

C MLE for logistic-quadratic-function-based approximation:

The cost and sigmoid function for this approximator is given below:

$$cost = -\frac{1}{N} \sum_{i=1}^N [y_i \ln(h(x_i, w)) + (1 - y_i) \ln(1 - h(x_i, w))]$$

where, the sigmoid function $h(x_i, w)$ is

$$h(x_i, w) = \frac{1}{1 + e^{-w^T z(x)}}$$

$$z(x) = [1, x_1, x_2, x_1^2, x_1 x_2, x_2^2]^T$$

i Using the 20 training samples:

After training the 20 samples data, the classifier was applied to the 10k validation datasets.




Figure 7: Classification of 10,000 validation samples based on 20 training samples

ii using the 200 training samples




Figure 8: Classification of 10,000 validation samples based on 200 training samples

iii using the 2000 training samples




Figure 9: Classification of 10,000 validation samples based on 2000 training samples

Training samples	minimum probability
20	0.223
200	0.171
2000	0.168

The following table shows the minimum probability between the logistic linear and quadratic classifiers.

Training samples	logistic-linear	logistic-quadratic
20	0.367	0.223
200	0.346	0.171
2000	0.343	0.168

Observation:

Similarly, in the quadratic approximation too, as the number of training samples increase the minimum probability error decreases. Furthermore, from the last table shown above, we can see that the quadratic approximation is better performing than the linear approximation, this is due to the overlapping Gaussian data distributions.

2. Given

$$y = c(x, w) + v; v \sim N(0, \sigma_{\text{noise}}); \gamma = [10^{-4} : 10^4]; z = [x^3, x^2, x, 1]^T$$

A **ML estimate**: The ML estimate is derived as follows and later implemented to estimate the coefficients of the cubic polynomial using 100 samples of data provided.

ML Estimator

$$w_{ML} = \underset{w}{\operatorname{argmax}} P(y|w, x)$$
$$-\log P(y|w, x) = -\log \prod_{n=1}^N P(y_n | w, x_n)$$
$$= -\sum_{n=1}^N \log P(y_n | w, x_n)$$

* In linear regression, the likelihood is Gaussian due to the Gaussian noise term $\epsilon \sim \mathcal{N}(0, \sigma^2)$.

$$\Rightarrow \log P(y|w, x, \sigma^2) = \log \left(\frac{1}{\sqrt{2\pi}\sigma^2} \exp \left(-\frac{(y_n - x_n^T w)^2}{2\sigma^2} \right) \right)$$
$$= -\frac{1}{2\sigma^2} (y_n - x_n^T w)^2 + k$$

\Rightarrow the loss function becomes;

$$L(w) = \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - x_n^T w)^2$$
$$= \frac{1}{2\sigma^2} (y - Xw)^T (y - Xw)$$
$$= \frac{1}{2\sigma^2} \|y - Xw\|^2$$

$$\frac{dL}{d\omega} = 0$$

$$\Rightarrow \frac{d}{d\omega} \left[\frac{1}{2\sigma^2} (y - X\omega)^T (y - X\omega) \right] = 0$$

$$\Rightarrow \frac{1}{2\sigma^2} (-2y^T X + 2\omega^T X^T X) = 0$$

$$\Rightarrow \omega^T X^T X = y^T X$$

$$\Rightarrow \omega^T = y^T X (X^T X)^{-1}$$

$$\Rightarrow \underline{\omega_{ML} = (X^T X)^{-1} X^T y}$$

Estimating Noise Variance

$$\log P(y|X, \omega, \sigma^2) = \sum_{n=1}^N \log \mathcal{N}(y_n | X_n \omega, \sigma^2)$$

$$= \sum_{n=1}^N \left(-\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\sigma^2) - \frac{1}{2\sigma^2} (y_n - X_n \omega)^2 \right)$$

$$\Rightarrow \frac{-2 \log P(y|X, \omega, \sigma^2)}{2\sigma^2} = \frac{-N}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{n=1}^N (y_n - X_n \omega)^2 = 0$$

$$\Rightarrow \underline{\underline{\sigma_{ML}^2 = \frac{1}{N} \sum_{n=1}^N (y_n - X_n \omega)^2}}$$

B **MAP estimate:** Similarly, the MAP estimate was derived for estimating the polynomial coefficients and later implemented using 100 training samples.

MAP estimator

$$p_{\text{prior}} = \mathcal{N}(0, \sigma^2 \mathbf{I})$$

$$P(w|x, y) \propto P(y|x, w) P(w)$$

* Negative log-likelihood $l(w)$:

$$l(w) = -\log P(w|x, y)$$

$$= -\log P(y|x, w) - \log P(w)$$

$$* \frac{\partial l(w)}{\partial w} = 0 = \frac{\partial}{\partial w} \left[\frac{1}{2\sigma^2} (y - Xw)^T (y - Xw) + \frac{1}{2\sigma^2} w^T w \right]$$

$$\Rightarrow w^T \left(\frac{1}{\sigma^2} X^T X + \frac{1}{\sigma^2} \mathbf{I} \right) - \frac{1}{\sigma^2} y^T X = 0$$

$$\Rightarrow \underline{\underline{w_{\text{MAP}}^T = \left(X^T X + \frac{\sigma^2}{\sigma^2} \mathbf{I} \right)^{-1} X^T y}}$$

After estimating the parameters /coefficients of the cubic polynomial, the model was implemented on 1000 validation samples and different metrics were used to test the performance of both estimators.

After estimating theta parameters for the approximator, I calculated the mean square error for both estimates.

The below graph shows the log-scale plot of mean square error of the validation dataset for ML and MAP estimate. In this case, the noise parameter v has a variance of 3.11 after training it separately, which has a closed form solution of

$$\frac{1}{N} \sum_{i=1}^N (y_i - xw)^2$$

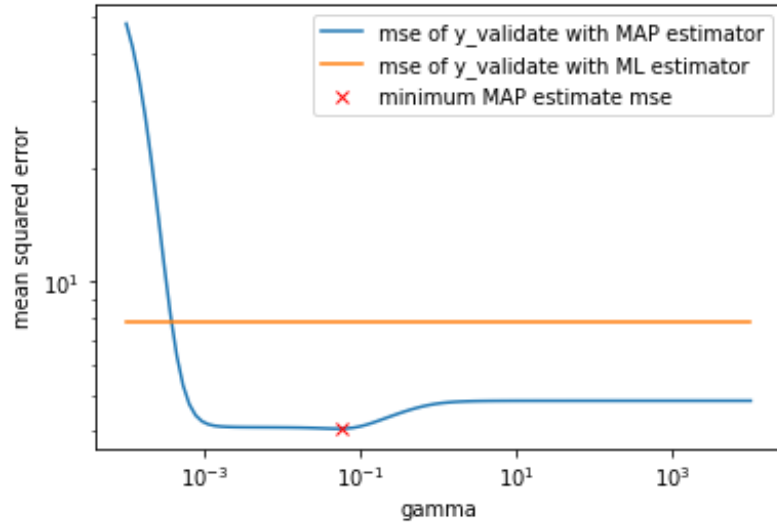


Figure 10: MSE of ML and MAP parameter estimations

The average MSE for using MAP estimator for different gamma values is 6.28, while MSE for ML was found to be 7.83. The minimum MSE for MAP estimator was 4.06 for $\gamma = 0.06$. From this, we can conclude that the optimal gamma value for the best fit is $\gamma = 0.06$. The MSE of MAP estimate after this optimal gamma value will be close to the ML estimate.

The l2 norm between the two θ parameter was also plotted:




Figure 11: l2 norm of $\theta_{ml} - \theta_{map}$

From the above figure, as the value of gamma increases, the l2 norm of $\theta_{ml} - \theta_{map}$ comes to a constant value (0.271), which implies for large value of gamma the MAP estimate becomes the same as the ML estimate. It can also be seen from the formula that, as the value of gamma increases the MAP estimate becomes ML estimate.

To see the effect of γ values on the training parameters (θ), I plotted the components of θ with γ values. For low values of γ , the θ components has same values, but as the value of γ increases, the parameter estimates starts to diverge.




Figure 12: θ_{map} components

$$3. \quad Z = [z_1, z_2, \dots, z_k]^T$$

$$\Theta = [\theta_1, \dots, \theta_k]^T$$

$$P(z_k = 1) = \theta_k$$

$D\{z_1, \dots, z_N\} \Rightarrow$ iid samples

a) ML estimator

$$\hat{\Theta}_{ML} = \arg \max_{\Theta} l(\Theta)$$

$$\begin{aligned} \text{where } l(\Theta) &= \ln P(Z_1 = z_1, \dots, Z_k = z_k; \theta_1, \dots, \theta_k) \\ &\stackrel{\text{iid}}{=} \ln \prod_{i=1}^N P(Z_i = z_i; \theta_1, \dots, \theta_k) \\ &= \sum_{i=1}^N \ln P(Z_i = z_i; \theta_1, \dots, \theta_k) \end{aligned}$$

$$\Rightarrow \hat{\Theta}_{ML} = \arg \max_{\Theta} \sum_{i=1}^N \ln P(Z_i = z_i | \Theta) \quad \text{s.t.}$$

$$\text{Constraints} \begin{cases} \Theta \geq 0 \\ \Theta^T \mathbf{1} = 1 \end{cases}$$

$$= \arg \min_{\Theta} -\frac{1}{N} \sum_{i=1}^N \ln \theta_{z_i} \quad \text{s.t.} \begin{cases} \Theta \geq 0 \\ \Theta^T \mathbf{1} = 1 \end{cases}$$

where $\mathbf{1} = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}$

Assume $\theta \geq 0$ will be inactive \Rightarrow
 $\theta_* > 0$

$$\Rightarrow \theta_* = \arg \min_{\theta} -\frac{1}{N} \sum_{i=1}^N \ln \theta_{z_i} \quad \text{s.t.} \quad \theta^T \mathbf{1} - 1 = 0$$

Lagrangian of the above optimization becomes

$$\mathcal{L}(\theta, \lambda) = -\frac{1}{N} \sum_{i=1}^N \ln \theta_{z_i} + \lambda (\theta^T \mathbf{1} - 1)$$

KKT conditions:

$$\nabla_{\theta} \mathcal{L}(\theta, \lambda) = 0$$

$$\Rightarrow -\frac{1}{N} \sum_{i=1}^N \frac{2}{2\theta_l} \ln \theta_{z_i} + \lambda \frac{2}{2\theta_l} (\theta^T \mathbf{1}) = 0$$

$$\Rightarrow \lambda - \frac{1}{N} \sum_{i=1}^M \frac{1}{\theta_l} \delta_{z_i l} = 0,$$

$$\text{where } \delta_{zl} = \begin{cases} 0 & \text{if } z \neq l \\ 1 & \text{if } z = l \end{cases}$$

$$\Rightarrow -\frac{1}{N\theta_e} \sum_{i=1}^N z_i l + \lambda = -\frac{N_e}{N\theta_e} - \lambda = 0$$

$$\Rightarrow \lambda = \frac{N_e}{N\theta_e}, \quad N_e \rightarrow \text{Number of samples with } z_i = l$$

$$N_1 + N_2 + \dots + N_K = N$$

$$\lambda \theta_e = \frac{N_e}{N}$$

$$\Rightarrow \lambda \theta_1 = \frac{N_1}{N}$$

$$\lambda \theta_2 = \frac{N_2}{N}$$

\vdots

$$\lambda \theta_K = \frac{N_K}{N}$$

Summing up the above:

$$\lambda (\theta_1 + \theta_2 + \dots + \theta_K) = \frac{N_1 + N_2 + \dots + N_K}{N} = 1$$

$$\Rightarrow \lambda = 1$$

$$\theta_e = \frac{N_e}{N}$$

b) MAP estimator

$$\text{Prior: } p(\theta|\alpha) = \frac{1}{B(\alpha)} \prod_{k=1}^K \theta_k^{\alpha_k-1}$$

$$B(\alpha) = \frac{\prod_{k=1}^K \Gamma(\alpha_k)}{\Gamma(\sum_{k=1}^K \alpha_k)}$$

$$\hat{\theta}_{\text{MAP}} = \underset{\theta}{\operatorname{argmax}} \ln p(\theta|D)$$

$$= \underset{\theta}{\operatorname{argmax}} \ln \frac{p(D|\theta) p(\theta)}{p(D)}$$

$$= \underset{\theta}{\operatorname{argmax}} \frac{1}{N} \sum_{i=1}^N \ln \theta_{\delta_i} + \ln \left(\frac{\prod_{k=1}^K \theta_k^{\alpha_k-1}}{B(\alpha)} \right)$$

$$\text{s.t.: } \theta \geq 0$$

$$\theta^T \mathbf{1} - 1 = 0$$

$$\begin{aligned} * \mathcal{L}(\theta, \lambda) = & \frac{1}{N} \sum_{i=1}^N \ln \theta_{\delta_i} + \sum_{k=1}^K (\alpha_k - 1) \ln \theta_k \\ & - \ln B(\alpha) - \lambda (\theta^T \mathbf{1} - 1) \end{aligned}$$

$$\frac{\partial \mathcal{L}}{\partial \theta_k} = 0 = \frac{N_k}{N\theta_k} + \sum_{k=1}^K \frac{(\alpha_k - 1)}{\theta_k} - \lambda$$

$$\Rightarrow \lambda \theta_k = \frac{N_k}{N} + \sum_k \alpha_k - K$$

$$= \frac{N_k}{N} + (\alpha^T \mathbf{1} - K)$$

\Rightarrow adding up the above expression for each k becomes?

$$\lambda (\theta_1 + \dots + \theta_K) = 1 + \alpha^T \mathbf{1} - K$$

$$\Rightarrow \lambda = (\alpha^T \mathbf{1} + 1 - K)$$

$$\Rightarrow \theta_k = \frac{\left(\frac{N_k}{N}\right) + (\alpha^T \mathbf{1} - K)}{(\alpha^T \mathbf{1} - K + 1)}$$

A Appendix

```
1
2 #Question 1
3 import matplotlib . pyplot as plt
4 import numpy as np
5 from scipy . stats import multivariate_normal as mvn
6 from collections import Iterable
7
8
9 features = 2
10 samples = [20,200,2000,10000]
11 mean_01=np.array([3,0])
12 mean_02=np.array([0,3])
13 mean_1=np.array([2,2])
14
15 cov_01=np.array([[2,0],[0,1]])
16 cov_02=np.array([[1,0],[0,2]])
17 cov_1=np.array([[1,0],[0,1]])
18
19 prior = [0.65, 0.35]
20 fpr = []
21 tpr = []
22
23 P_error = []
24 gamma_list = []
25
26 #part one
27
28 def data(n):
29     label = np.zeros ((3, n ))
30     label [0 , :] = (np.random.uniform (0,1,n)>= prior[0]).astype (int)
31     for i in range(n):
32         p=np.random.uniform(0,1)
33         if p>=prior[0]:
34             label[0,i]=1
35         else:
36             label[0,i]=0
37     x = np.zeros ((features , n))
38     for index in range (n) :
39         if label [0,index] == 0:
40             w=np.random.uniform (0,1)
41             if w>=0.5:
42                 x[:,index] = mvn(mean_01.reshape(2,),cov_01).rvs(1)
43             else:
44                 x[:,index] = mvn(mean_02.reshape(2,),cov_02).rvs(1)
45         else:
46             x[:,index] = mvn(mean_1.reshape(2,),cov_1).rvs(1)
47     return x,label
48
49 Data_train_20,label_20=data(samples[0])
50 Data_train_200,label_200=data(samples[1])
51 Data_train_2000,label_2000=data(samples[2])
52 Data_validate_10000,label_10000=data(samples[3])
53
54 logValpdf0=np.zeros(samples[3])
55
56 for i in range(samples[3]):
```

```

57 pdf01= np.log(mvn.pdf(Data_validate_10000[:,i].T,mean = mean_01,cov =
    cov_01))
58 pdf02= np.log(mvn.pdf(Data_validate_10000[:,i].T,mean = mean_02,cov =
    cov_02))
59 logValpdf0[i] = 0.5*pdf01+0.5*pdf02
60
61 logValpdf1= np.log(mvn.pdf(Data_validate_10000.T,mean=mean_1, cov =
    cov_1))
62 discriminant_score = logValpdf1 - logValpdf0
63 class0_count = float(list(label_10000[0,:]).count(0)) # number of
    samples for class 0
64 class1_count = float(list(label_10000[0,:]).count(1)) # number of
    samples for class 1
65
66 tau=(sorted(discriminant_score[np.array(discriminant_score[:]).astype(
    float) >=0]))
67 mid_tau=np.array([tau[0]-100,(tau[0:(len(tau)-1)]+(np.diff(tau))/2).
    tolist(),tau[len(tau)-1]+100])
68
69 def flatten(lis):
70     for item in lis:
71         if isinstance(item, Iterable) and not isinstance(item, str):
72             for x in flatten(item):
73                 yield x
74         else:
75             yield item
76 mid_tau=list(flatten(mid_tau.tolist()))
77
78
79 for gamma in mid_tau:#gamma_list
80     label_10000[1,:] = (discriminant_score>=gamma)#.astype(int)#np.log(
        gamma)
81
82 x01 = [i for i in range(label_10000.shape[1]) if (label_10000[1,i] == 1
    and label_10000[0,i] == 0)]
83 x11 = [i for i in range(label_10000.shape[1]) if (label_10000[1,i] == 1
    and label_10000[0,i] == 1)]
84 tpr.append(len(x11)/class1_count)
85 fpr.append(len(x01)/class0_count)
86 P_error.append( (len(x01)/class0_count)*prior[0] + (1-len(x11)/
    class1_count)*prior[1])
87
88
89 minimum_Perror=min(P_error)
90 min_idx=np.argmin(P_error)
91 fpr_theory=[]
92 tpr_theory=[]
93 p_thry=[]
94
95 label_10000[2,:] = (discriminant_score>= np.log(prior[0]/prior[1])).
    astype(int)
96 x00t = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] == 0
    and label_10000[2,i] == 0)]
97 x01t = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] == 0
    and label_10000[2,i] == 1)]#fp
98 x10t = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] == 1
    and label_10000[2,i] == 0)]#fN
99 x11t = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] == 1

```

```

        and label_10000[2,i] == 1))
100 fpr_theory . append (len( x01t ) / class0_count )
101 tpr_theory . append (len( x11t ) / class1_count )
102
103 p_thry.append( (len(x01t)/class0_count)*prior[0] + (1-len(x11t)/
        class1_count)*prior[1])
104
105 print('Optimal threshold {}'.format(mid_tau[min_idx]))#change here
106 print('TPR at minimum probability :{}'.format ( tpr [ min_idx ]))
107 print('FPR at minimum probability :{}'.format ( fpr [ min_idx ]))
108 print('Minimum Probability Error:{}'.format(minimum_Perror))
109 print('Minimum Theoretical Probability Error:{}'.format(min(p_thry)))
110
111 #plot ROC
112
113 plt.plot(fpr,tpr,color = 'red' )
114 plt.plot(fpr[min_idx],tpr[min_idx], 'o',color='k')
115 plt.xlabel('FPR')
116 plt.ylabel('TPR')
117 plt.title('ROC Curve')
118 plt.legend(['ROC','Experimental min Error'])
119 plt.show()
120
121
122
123 #plot the predicted labels
124 plt.plot(Data_validate_10000[0,x00t],Data_validate_10000[1,x00t], '.',
        color = 'g', markersize = 6)
125 plt.plot(Data_validate_10000[0,x01t],Data_validate_10000[1,x01t], '.',
        color = 'r', markersize = 6)
126 plt.plot(Data_validate_10000[0,x11t],Data_validate_10000[1,x11t], '+',
        color = 'g', markersize = 6)
127 plt.plot(Data_validate_10000[0,x10t],Data_validate_10000[1,x10t], '+',
        color = 'r', markersize = 6)
128 plt.legend(['class 0 correctly classified','class 0 wrongly classified',
        'class 1 correctly classified','class 1 wrongly classified'])
129 plt.xlabel("x1")
130 plt.ylabel("x2")
131
132 horizontalGrid = np.linspace(np.floor(min(Data_validate_10000[0,:])),np.
        ceil(max(Data_validate_10000[0,:])),100)
133 verticalGrid = np.linspace(np.floor(min(Data_validate_10000[1,:])),np.
        ceil(max(Data_validate_10000[1,:])),100)
134
135 dsg = np.zeros((100,100))
136 a = np.array(np.meshgrid(horizontalGrid,verticalGrid))
137 for i in range(100):
138     for j in range(100):
139         ww=np.random.uniform(0,1)
140         p = mvn.pdf(np.array([a[0][i][j], a[1][i][j]]),mean=mean_1, cov =
        cov_1)
141         q1=mvn.pdf(np.array([a[0][i][j], a[1][i][j]]),mean=mean_01, cov =
        cov_01)
142         q2=mvn.pdf(np.array([a[0][i][j], a[1][i][j]]),mean=mean_02, cov =
        cov_02)
143         q=0.5*q1+0.5*q2
144         dsg[i][j] = np.log(q) - np.log(p) - np.log(prior[0]/prior[1])
145 plt.contour(a[0],a[1],dsg,levels=[0])

```

```

146 plt.show()
147
148
149 x0 = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] == 0)]
150 x1 = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] == 1 )
151 ]
152 plt.plot(Data_validate_10000[0,x0],Data_validate_10000[1,x0],'+')
153 plt.plot(Data_validate_10000[0,x1],Data_validate_10000[1,x1],'.')
154 plt.xlabel('x1')
155 plt.ylabel('x2')
156 plt.title("Actual Data distribution")
157 plt.legend(['Class 0','Class 1'])
158 plt.show()
159
160 #part 2
161 def linear_reg(xdata,alpha,max_iteration,ydata,xtest):
162     z=np.vstack((np.ones(xdata.shape[1]),xdata))
163     w=np.zeros((3,1))
164
165     for i in range(max_iteration):
166         h=1/(1+np.exp(-(np.dot(w.T,z))))
167         grad=(1/float(z.shape[1])) * np.dot(z,(h-ydata[0]).T)
168         w=w-alpha*grad
169     z=np.vstack((np.ones(xtest.shape[1]),xtest))
170     decisions=np.zeros((1,xtest.shape[1]))
171     h=1/(1+np.exp(-(np.dot(w.T,z))))
172     decisions[0,:]=(h[0,:]>=0.5).astype(int)
173     return w,decisions
174 def quadratic_reg(xdata,alpha,max_iteration,ydata,xtest):
175
176     z = np.vstack((np.ones(xdata.shape[1]),xdata[0],xdata[1],xdata[0]*
177         xdata[0],xdata[0]*xdata[1], xdata[1]*xdata[1]))
178     w = np.zeros((6,1))
179     for i in range(max_iteration):
180         h=1/(1+np.exp(-(np.dot(w.T,z))))
181         grad=(1/float(z.shape[1])) * np.dot(z,(h-ydata[0]).T)
182         w=w-alpha*grad
183     z = np.vstack((np.ones(xtest.shape[1]),xtest[0],xtest[1],xtest[0]*
184         xtest[0],xtest[0]*xtest[1], xtest[1]*xtest[1]))
185     decisions = np.zeros((1,xtest.shape[1]))
186     h=1/(1+np.exp(-(np.dot(w.T,z))))
187     decisions[0,:]=(h[0,:]>=0.5).astype(int)
188     return w,decisions
189
190 alpha=0.05
191 max_iter=2000
192
193 w_20,decisions_20=linear_reg(Data_train_20,alpha,max_iter,label_20,
194     Data_validate_10000)
195 w_200,decisions_200=linear_reg(Data_train_200,alpha,max_iter,label_200,
196     Data_validate_10000)
197 w_2000,decisions_2000=linear_reg(Data_train_2000,alpha,max_iter,
198     label_2000,Data_validate_10000)
199
200 w_20q,decisions_20q=quadratic_reg(Data_train_20,alpha,max_iter,label_20,
201     Data_validate_10000)

```

```

197 w_200q,decisions_200q=quadratic_reg(Data_train_200,alpha,max_iter,
    label_200,Data_validate_10000)
198 w_2000q,decisions_2000q=quadratic_reg(Data_train_2000,alpha,max_iter,
    label_2000,Data_validate_10000)
199
200 def contour_plot(typ,w_par):
201     horizontalGrid = np.linspace(np.floor(min(Data_validate_10000[0,:])),
        np.ceil(max(Data_validate_10000[0,:])),100)
202     verticalGrid = np.linspace(np.floor(min(Data_validate_10000[1,:])),np.
        ceil(max(Data_validate_10000[1,:])),100)
203     dsg = np.zeros((100,100))
204     a = np.array(np.meshgrid(horizontalGrid,verticalGrid))
205     for i in range(100):
206         for j in range(100):
207             x1 = a[0][i][j]
208             x2 = a[1][i][j]
209             if typ==0:
210                 z = np.c_[1,x1,x2].T
211             else:
212                 z = np.c_[1,x1,x2,pow(x1,2),x1*x2,pow(x2,2)].T
213             dsg[i][j] = np.sum(np.dot(w_par.T,z))
214     plt.contour(a[0],a[1],dsg,levels=[0])
215     plt.show()
216
217 def plots(Data_train,label_train,decision,w_val,typ):
218     P_err=[]
219
220     #actual data plots
221     x0 = [i for i in range(label_train.shape[1]) if (label_train[0,i] ==
        0)]
222     x1 = [i for i in range(label_train.shape[1]) if (label_train[0,i] == 1
        )]
223
224     plt.plot(Data_train[0,x0],Data_train[1,x0],'+')
225     plt.plot(Data_train[0,x1],Data_train[1,x1],'.')
226     plt.xlabel('x1')
227     plt.ylabel('x2')
228     plt.title("Actual Data distribution")
229     plt.legend(['Class 0','Class 1'])
230     if typ==0:#linear
231         contour_plot(0,w_val)
232     else:
233         contour_plot(1,w_val)
234
235     #plot the classified data
236
237     x00 = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] ==
        0 and decision[0,i] == 0)]
238     x01 = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] ==
        0 and decision[0,i] == 1)]#FP
239     x10 = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] ==
        1 and decision[0,i] == 0)]
240     x11 = [i for i in range(label_10000.shape[1]) if (label_10000[0,i] ==
        1 and decision[0,i] == 1)]#TP
241     P_err.append((len(x01)/class0_count)*prior[0] + (1-len(x11)/
        class1_count)*prior[1])
242     print(min(P_err))
243

```

```

244 plt.plot(Data_validate_10000[0,x00],Data_validate_10000[1,x00],'.',
    color = 'g', markersize = 6)
245 plt.plot(Data_validate_10000[0,x01],Data_validate_10000[1,x01],'.',
    color = 'r', markersize = 6)
246 plt.plot(Data_validate_10000[0,x11],Data_validate_10000[1,x11],'+',
    color = 'g', markersize = 6)
247 plt.plot(Data_validate_10000[0,x10],Data_validate_10000[1,x10],'+',
    color = 'r', markersize = 6)
248 plt.legend(['class 0 correctly classified','class 0 wrongly classified',
    'class 1 correctly classified','class 1 wrongly classified'])
249 plt.xlabel("x1")
250 plt.ylabel("x2")
251
252 if typ==0:#linear
253     contour_plot(0,w_val)
254 else:
255     contour_plot(1,w_val)
256
257
258 plots(Data_train_20,label_20,decisions_20,w_20,0)
259 plots(Data_train_20,label_20,decisions_20q,w_20q,1)
260 plots(Data_train_200,label_200,decisions_200,w_200,0)
261 plots(Data_train_200,label_200,decisions_200q,w_200q,1)
262 plots(Data_train_2000,label_2000,decisions_2000,w_2000,0)
263 plots(Data_train_2000,label_2000,decisions_2000q,w_2000q,1)

```

Listing 1: Question 1

```

1 #Question 2
2 import numpy as np
3 import pandas as pd
4 from numpy.matlib import repmat
5 import matplotlib.pyplot as plt
6
7
8 def generateData(N):
9     gmmParameters = {}
10    gmmParameters['priors'] = [.3,.4,.3] # priors should be a row vector
11    gmmParameters['meanVectors'] = np.array([[ -10, 0, 10], [0, 0, 0],
12    [10, 0, -10]])
13    gmmParameters['covMatrices'] = np.zeros((3, 3, 3))
14    gmmParameters['covMatrices'][:, :, 0] = np.array([[1, 0, -3], [0, 1,
15    0], [-3, 0, 15]])
16    gmmParameters['covMatrices'][:, :, 1] = np.array([[8, 0, 0], [0, .5,
17    0], [0, 0, .5]])
18    gmmParameters['covMatrices'][:, :, 2] = np.array([[1, 0, -3], [0, 1,
19    0], [-3, 0, 15]])
20    x, labels = generateDataFromGMM(N, gmmParameters)
21    return x
22
23 def generateDataFromGMM(N, gmmParameters):
24     # Generates N vector samples from the specified mixture of Gaussians
25     # Returns samples and their component labels
26     # Data dimensionality is determined by the size of mu/Sigma
27     parameters
28     priors = gmmParameters['priors'] # priors should be a row vector
29     meanVectors = gmmParameters['meanVectors']
30     covMatrices = gmmParameters['covMatrices']
31     n = meanVectors.shape[0] # Data dimensionality

```



```

27     C = len(priors) # Number of components
28     x = np.zeros((n,N))
29     labels = np.zeros((1,N))
30     # Decide randomly which samples will come from each component
31     u = np.random.random((1,N))
32     thresholds = np.zeros((1,C+1))
33     thresholds[:,0:C] = np.cumsum(priors)
34     thresholds[:,C] = 1
35     for l in range(C):
36         indl = np.where(u <= float(thresholds[:,l]))
37         Nl = len(indl[1])
38         labels[indl] = (l+1)*1
39         u[indl] = 1.1
40         x[:,indl[1]] = np.transpose(np.random.multivariate_normal(
meanVectors[:,l], covMatrices[:, :, l], Nl))
41
42     return x, labels
43
44 def plot3(a,b,c,mark="o",col="b"):
45     from matplotlib import pyplot
46     import pylab
47     from mpl_toolkits.mplot3d import Axes3D
48     pylab.ion()
49     fig = pylab.figure()
50     ax = Axes3D(fig)
51     ax.scatter(a, b, c, marker=mark, color=col)
52     ax.set_xlabel("x1")
53     ax.set_ylabel("x2")
54     ax.set_zlabel("y")
55     ax.set_title('Training Dataset')
56
57 def mse(x,y):
58     error=0
59     for i in range(x.shape[0]):
60         error=error+pow(x[i]-y[i],2)
61
62     return (1/x.shape[0])*error
63
64 Ntrain = 100
65 data1 = generateData(Ntrain)
66 #plot3(data1[0,:],data1[1,:],data1[2,:])#plot the training data
67 xTrain= data1[0:2,:]
68 yTrain = data1[2,:]
69
70 Ntrain = 1000
71 data2 = generateData(Ntrain)
72 #plot3(data2[0,:],data2[1,:],data2[2,:])#plot the validation data
73 xValidate = data2[0:2,:]
74 yValidate = data2[2,:]
75 sigm=0.1
76
77 z=np.vstack((pow(xTrain,3),pow(xTrain,2),pow(xTrain,1),pow(xTrain,0)))
#design matrix
78
79 gamma_list=[]
80
81 for i in np.logspace(-4,4,100):
82     gamma_list.append(i)

```



```

83
84 zzt=np.zeros((8,8))
85 b=np.zeros(8).reshape(8,1)
86 for j in range(z.shape[1]):
87     temp=z[:,j].reshape(8,1)
88     zzt=zzt+np.matmul(temp,temp.T)
89     b=b+temp*(yTrain[j])
90 theta_map_list=[]
91
92 theta_ml=np.matmul(np.linalg.inv(zzt+0.0001*np.random.normal(0,sigm
    ,(8,8))),b)# I added the small error to make the matrix non-singular
93
94 sigmaa=0
95 for i in range(xTrain.shape[1]):
96     temp=np.matmul(theta_ml[0:2,0].reshape(1,2), pow(xTrain[:,i],3)) + np.
        matmul(theta_ml[2:4,0].reshape(1,2), pow(xTrain[:,i],2))+np.matmul(
            theta_ml[4:6,0].reshape(1,2), pow(xTrain[:,i],1))\
97 +np.matmul(theta_ml[6:8,0].reshape(1,2), pow(xTrain[:,i],0))
98     sigmaa=sigmaa+pow((yTrain[i]-temp),2)
99
100 sigmaa=(1/xTrain.shape[1])*sigmaa
101
102
103 for gamma in gamma_list:
104     A=zzt+(sigmaa/pow(gamma,2))*np.eye(8)
105     theta_map=np.matmul(np.linalg.inv(A),b)
106
107     theta_map_list.append(theta_map)
108
109 v = np.random.normal(0,np.sqrt(sigmaa),1000)
110 y_validate_ml=np.matmul(theta_ml[0:2,0].reshape(1,2),pow(xValidate,3)) +
        np.matmul(theta_ml[2:4,0].reshape(1,2),pow(xValidate,2))\
111 +np.matmul(theta_ml[4:6,0].reshape(1,2),pow(xValidate,1))+np.matmul(
        theta_ml[6:8,0].reshape(1,2),pow(xValidate,0))+v
112
113
114 y_validate_map=[]
115 l2_norm=[]
116 l2a,l2b,l2c,l2d=[[],[],[],[]]
117 y_train_map=[]
118
119 for i in theta_map_list:
120     y_val=np.matmul(i[0:2,0].reshape(1,2),pow(xValidate,3)) + np.matmul(i
        [2:4,0].reshape(1,2),pow(xValidate,2))\
121 +np.matmul(i[4:6,0].reshape(1,2),pow(xValidate,1))+np.matmul(i[6:8,0].
        reshape(1,2),pow(xValidate,0))
122
123     y_validate_map.append(y_val)
124     l2_norm.append(pow(np.linalg.norm(theta_ml-i),2))
125     y_train_map.append(np.matmul(i[0:2,0].reshape(1,2),pow(xTrain,3)) + np
        .matmul(i[2:4,0].reshape(1,2),pow(xTrain,2))\
126 +np.matmul(i[4:6,0].reshape(1,2),pow(xTrain,1))+np.matmul(i[6:8,0].
        reshape(1,2),pow(xTrain,0)))
127
128
129 mse_list_map=[]
130 mse_list_map_train=[]
131 for y in y_validate_map:

```

```

132     rr=mse(y.T,yValidate.T)
133     mse_list_map.append(rr)
134
135 a=[]
136 b=[]
137 c=[]
138 d=[]
139 for i in theta_map_list:
140     a.append(i[0:2,0])
141     b.append(i[2:4,0])
142     c.append(i[4:6,0])
143     d.append(i[6:8,0])
144
145 for xx in y_train_map:
146     mse_list_map_train.append(mse(xx.T,yTrain.T))
147
148 mse_list_ml=mse(yValidate.T,y_validate_ml.T)
149
150 from numpy.lib.function_base import average
151 print('mse for using the ml estimator:{}'.format(mse_list_ml))
152 print('average mse for using the map estimator for different gamma
    values:{}'.format(average(mse_list_map)))
153 print('minimum map mse:{}'.format(min(mse_list_map)))
154 print(gamma_list[np.argmin(mse_list_map)])
155
156 plt.plot(gamma_list,l2_norm,label='l2 norm')
157 plt.plot(gamma_list,np.zeros(len(gamma_list)), '--',label='reference')
158 plt.xscale(value='log')
159 plt.yscale(value='log')
160 plt.xlabel('gamma values')
161 plt.ylabel('l2 norm between theta map and ml')
162 plt.legend()
163 plt.show()
164
165 plt.plot(gamma_list,mse_list_map,label='mse of y_validate with MAP
    estimator')
166 #plt.plot(gamma_list,mse_list_map_train,label='mse of y_train map with
    gamma values')
167 plt.plot(gamma_list,[mse_list_ml for i in range(len(gamma_list))],label=
    'mse of y_validate with ML estimator')
168 plt.plot(gamma_list[np.argmin(mse_list_map)],min(mse_list_map),'xr',
    label='minimum MAP estimate mse')
169 plt.xscale(value='log')
170 plt.yscale(value='log')
171 plt.xlabel('gamma')
172 plt.ylabel('mean squared error')
173 plt.legend()
174 plt.show()
175
176
177 plt.plot(gamma_list,a)
178 plt.plot(gamma_list,b)
179 plt.plot(gamma_list,c)
180 plt.plot(gamma_list,d)
181 plt.plot(gamma_list,np.zeros(len(gamma_list)), '--',label='reference')
182 plt.xscale(value='log')
183 #plt.yscale(value='log')
184 plt.xlabel('gamma')

```

```
185 plt.ylabel(r'MAP estimated parameters ($\theta$)')
186 plt.legend([r'$\theta_{00}$',r'$\theta_{01}$',r'$\theta_{10}$',r'$\theta_{11}$',r'$\theta_{20}$',r'$\theta_{21}$',r'$\theta_{30}$',r'$\theta_{31}$'])
```

Listing 2: Question 2