

EECE 5644: Machine Learning Project 4 Report

1. Given the following data:

$$Priors = [0.5, 0.5]$$

$$number\ of\ samples = [1000, 10000]$$

$$\mathbf{x} = r_l \begin{bmatrix} \cos(\theta) \\ \sin(\theta) \end{bmatrix} + \mathbf{n}$$

where, $\theta \sim \text{uniform}[-\pi, \pi]$, $\mathbf{n} \sim \mathcal{N}(0, \sigma^2 I)$, $r_{-1} = 2, r_{+1} = 4, \sigma = 1$

Actual Data Distribution:




Figure 1: Actual training data distribution




Figure 2: Actual test data distribution

I MLP

Implementation:

After generating the samples based on the above expression, and Gaussian parameters (zero mean and identity covariance), I used the training dataset (1000 samples) for choosing the hyper-parameters and later apply MLP on the test dataset.

10-fold cross validation was used in order to select the optimal number of perceptrons, optimal activation function type for the hidden layer, optimal number of epochs and optimizer type. To find the optimal number of epochs, I used a log-space function with base 5 and maximum limit of 625. Activation functions such as *ReLU*, *sigmoid*, and *tanh* were compared with their minimum probability of error. After choosing the right epoch number and activation function type, I apply the same procedure to find the optimal number of perceptrons with the selected epoch number and activation function type, and the perceptron with the minimum error was selected. Similarly, the optimizers were compared while I was choosing the number of perceptrons. Since the number of class is 2, I used a *Sigmoid* activation function at the output layer. Using the hyper-parameters I got from the above procedure, the classifier was applied on the test dataset (10K), and the confusion matrix along with the error percentage and classification performance (with a threshold value of 0.5) plot with decision boundary super-imposed were reported.

a Number of epochs

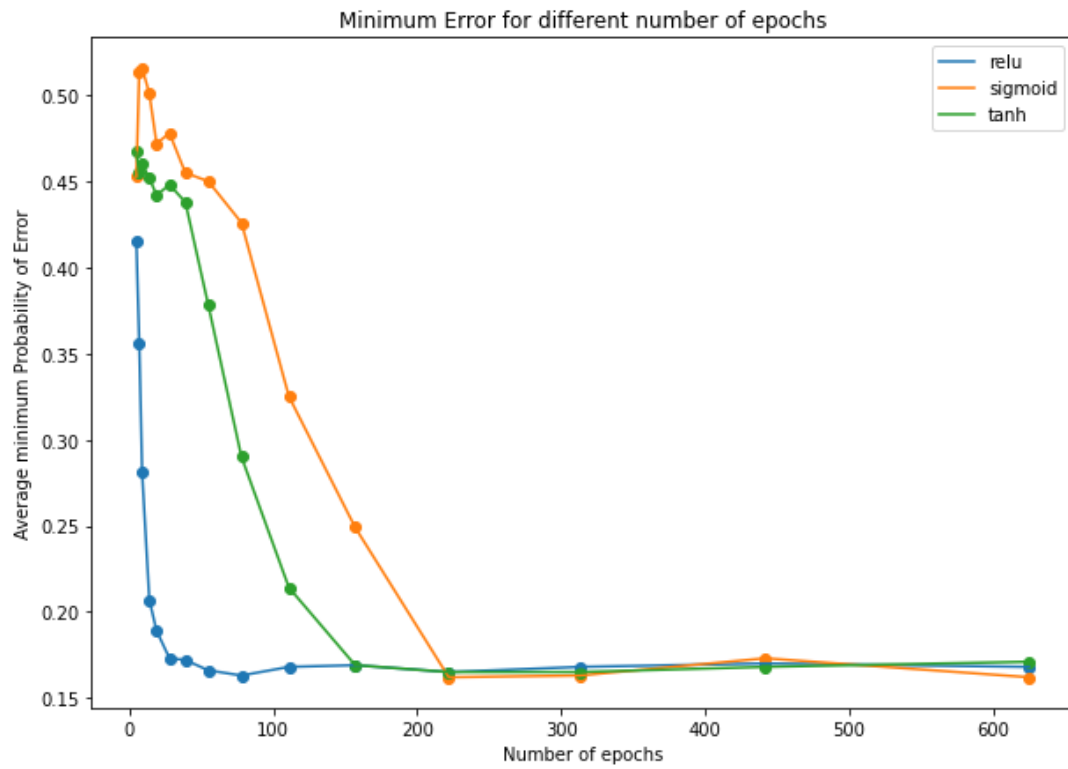


Figure 3: Minimum Error for different number of epochs

b Number of Perceptrons

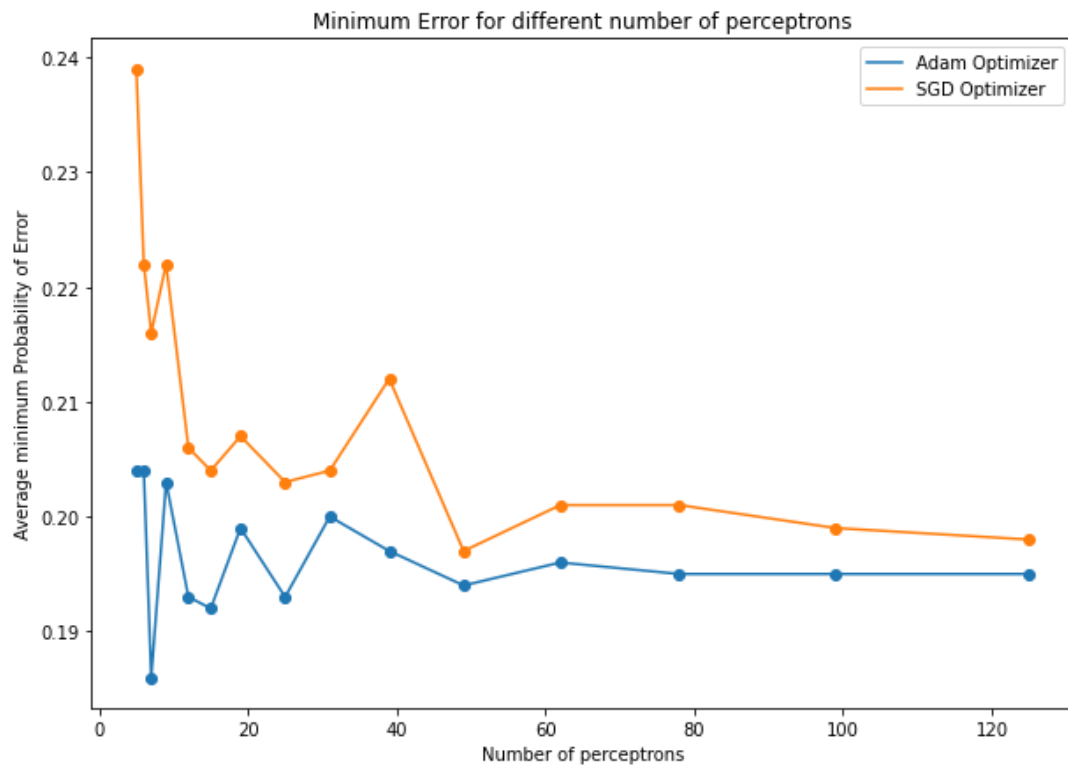


Figure 4: Minimum Error for different number of neurons under ReLU activation function

Minimum probability of error for Adam:0.186
Minimum probability of error for SGD:0.197

Based on the above plots and probability error, the following table is prepared and shows the selected parameter values for the test dataset.

Number of epochs	100
Batch size	32
Number of perceptron	7
Optimizer	Adam
Activation function for hidden layer	ReLU
Activation function for output layer	Sigmoid
loss function	binary cross_entropy

c Confusion Matrix

The test dataset (10K), was tested based on the above parameters (in the table) given, and the confusion matrix along with the minimum probability of error is reported below. For the confusion matrix, rows are actual values while column indicates predicted values.

Confusion matrix:

Normalized Confusion matrix:

$$\begin{bmatrix} 4066 & 950 \\ 681 & 4303 \end{bmatrix} \quad \begin{bmatrix} 0.81 & 0.19 \\ 0.14 & 0.86 \end{bmatrix}$$

Minimum probability of error: 0.163
Accuracy: 83.7%
cross_entropy loss:0.402

d Classification Performance:




Figure 5: MLP Classification performance

II SVM

Implementation: The same data as above was used for this classifier, and 10-fold cross validation was performed on the training dataset to choose the hyper-parameters(C, γ) from the following combinations, $C=[1000,100,10,1.0,0.1,0.01]$ and $\gamma=[0.01,0.1,1.0,10,100,1000]$. For choosing the optimal hyper-parameters, I used *GridSearchCV* function from sklearn library, and a combination of $C=100$ and $\gamma=0.01$ has highest accuracy (84.1%) and applied on the test dataset. A classification performance (with threshold value of 0.5) was plotted at the end, which shows the correct and incorrect labels with the decision boundary superimposed.

a Hyper-parameter Selection

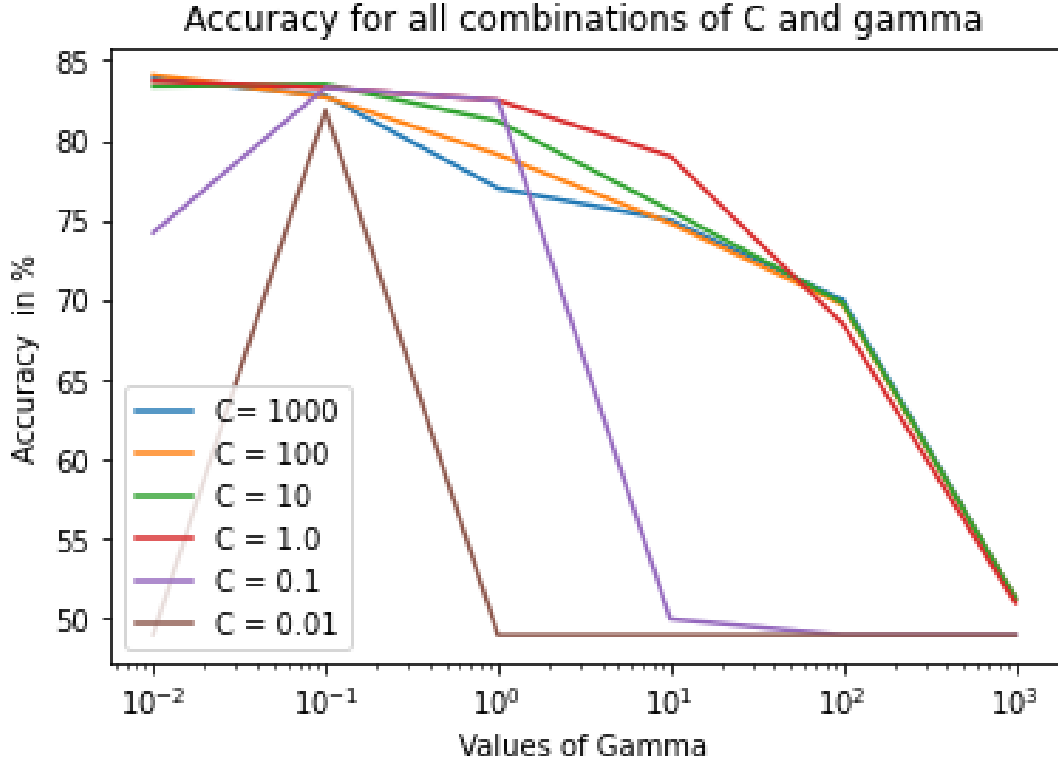


Figure 6: Accuracy for different combinations of hyper-parameters

Based on the above plots the following table is prepared and shows the selected parameter values (with the highest accuracy) for the test dataset.

C	100
γ	0.01
kernel	rbf

b Confusion Matrix

The test dataset (10K), was tested based on the above parameters (in the table) given, and the confusion matrix along with the minimum probability of error is reported below. For the confusion matrix, rows are actual values while column indicates predicted values.

Confusion matrix:

Normalized Confusion matrix:

$$\begin{bmatrix} 4202 & 841 \\ 822 & 4135 \end{bmatrix} \quad \begin{bmatrix} 0.83 & 0.17 \\ 0.16 & 0.83 \end{bmatrix}$$

Minimum probability of error: 0.166

Accuracy: 83.4%

c Classification Performance:

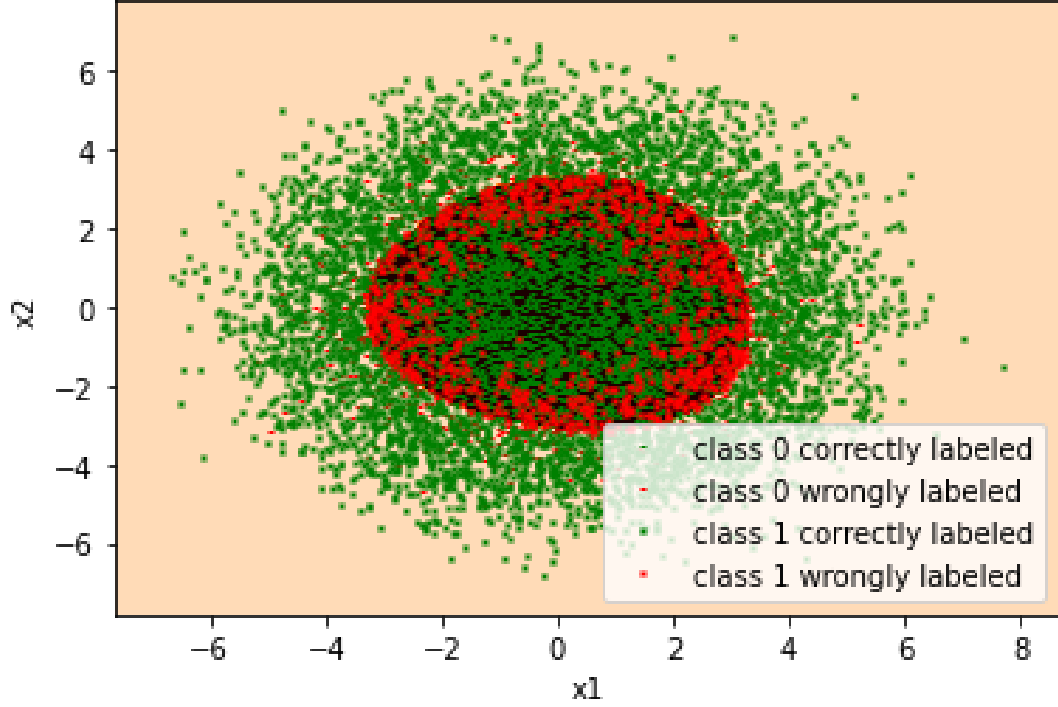


Figure 7: SVM Classification performance

III Conclusion

A comparison between the two classifiers is summarized with the below table.

Classifier	Training Accuracy (%)	Test Accuracy(%)
MLP	81.4	83.7
SVM	84.1	83.4

From the above table, we can see that the training accuracy for the SVM classifier is higher than the MLP, while for the test dataset the classification accuracy is almost equal. As the number of samples increases (from 1000 to 10k), there will be more samples in the decision boundary and as a result the SVM performs poor as the samples increase. While for MLP, the accuracy increases as samples increases because it will capture the true distribution of data and this will minimize both bias and variance.

2. GMM

Implementation:

I used three color images (deer,street and Penguin), from the Berkeley Segmentation Dataset and generate 5 features using the row index, column index ,red,green and blues pixel values. After forming the feature vector and scaled it to $[0,1]$ using *min_maxscale* function from sklearn library, I used the built in function *GaussianMixture* from sklearn library to estimate the parameters for each Gaussian component. Expectation maximization (EM) algorithm was used in-order to maximize the expected value of the log likelihood function. To find the best Gaussian component, 10 fold cross validation was used. The performance was measured by sklearn library function *sklearn.model_selection.cross_val_score*, which return log likelihood score of data. This score was used to find the best model order (which has maximum score).

I Image 1 (Deer)

a Model order selection




Figure 8: log-likelihood score for different clusters (model order)

As it can be seen from the above figure, a model order of 5 has the highest log likelihood score and was selected for the Gaussian component while segmenting the input image.

b Results



(a) Original image



(b) Segmented image

II Image 2(street)

a Model order selection

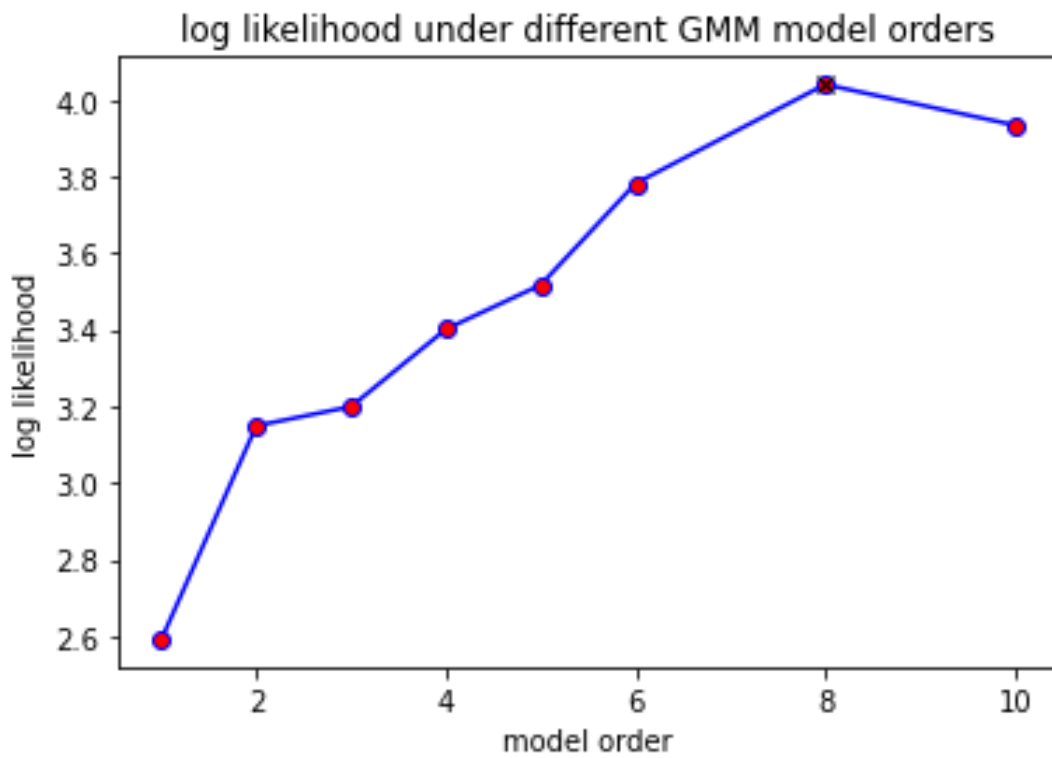



Figure 10: log-likelihood score for different clusters (model order)

As it can be seen from the above figure, a model order of 8 has the highest log likelihood score and was selected for the Gaussian component while segmenting the input image.

b Results



(a) Original image



(b) Segmented image

III Image 3 (Penguin)

a Model order selection

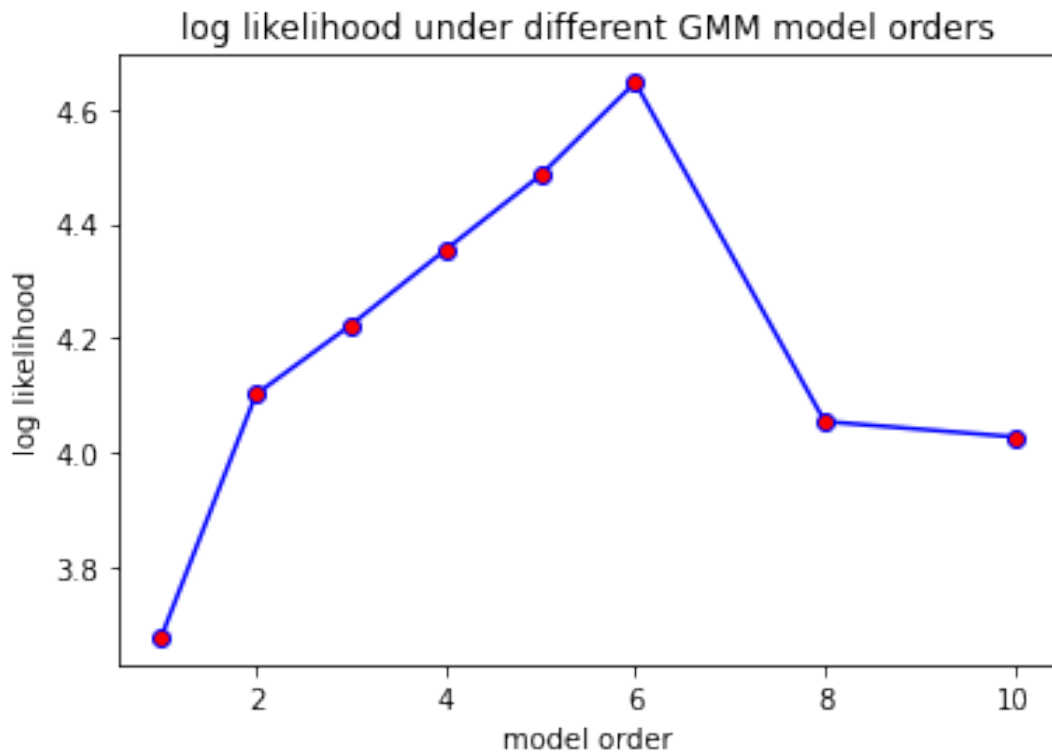



Figure 12: log-likelihood score for different clusters (model order)

As it can be seen from the above figure, a model order of 6 has the highest log likelihood score and was selected for the Gaussian component while segmenting the input image.

b Results



(a) Original image



(b) Segmented image

A Appendix

```
1
2
3 # Question 1
4 import numpy as np
5 from scipy . stats import multivariate_normal as mvn
6 import tensorflow as tf
7 from sklearn.metrics import confusion_matrix
8 import random
9 from tensorflow.keras.utils import to_categorical
10 import matplotlib.pyplot as plt
11 from tensorflow.keras.layers import Dense
12 import matplotlib.pyplot as plt
13 import warnings
14 warnings.filterwarnings("ignore", category=RuntimeWarning)
15 import math
16
17 from sklearn.svm import SVC
18 from sklearn.metrics import accuracy_score
19 from sklearn.model_selection import GridSearchCV, KFold
20 from sklearn.metrics import classification_report
21
22
23 def data(n):
24     label = np.zeros((1,n))
25     for i in range(n):
26         label[0,i]= np.random.choice([0,1],p=prior) # equal prior
27
28     x = np.zeros ((features , n))
29
30     for index in range (n) :
31         theta=np.random.uniform(theta_margin[0],theta_margin[1])
32         temp=np.array([np.cos(theta),np.sin(theta)])
33
34         if label [0,index] == 0:#sample from class -1
35             x[:,index] = 2*(temp)+mvn(m1_n,c1_n).rvs(1)
36         else:# from class +1
37             x[:,index] = 4*(temp)+mvn(m2_n,c2_n).rvs(1)
38
39     return x,label
40
41
42 def split_data(data, folds):
43     data_split = []
44     random.seed(1)
45     data_temp = list(data)
46     fold_size = int(len(data) / folds)
47     for i in range(folds):
48         fold = list()
49         while len(fold) < fold_size:
50             rand_index = random.randrange(len(data_temp))
51             fold.append(data_temp.pop(rand_index))
52         data_split.append(fold)
53     return data_split
54
55 def test_train_split(folds,i):
56     test=np.array(folds.pop(i))
```

```

57 train=np.array(folds)
58
59 xtest=test[:,0:2]
60 ytest=test[:,2]
61 xtrain=train[:, :,0:2].reshape(-1,2)
62 ytrain=train[:, :,2].reshape(-1)
63
64 return xtrain,ytrain,xtest,ytest
65
66 def MLP(sample_type):
67
68     error_mat=np.zeros((len(percept_list),len(activations)))
69     accuracy_mat=np.zeros((len(percept_list),len(activations)))
70
71     data_concat=np.hstack((X_train,Y_train))
72
73     for activ in range(len(activations)):
74
75         for idx,percpt in enumerate(percept_list):
76             err_fold=[]
77             acc_fold=[]
78
79             for i in range(num_folds):
80                 X_train_,Y_train_,X_test_,Y_test_=test_train_split(split_data(
data_concat,num_folds),i)
81                 input_shape = (features,)
82                 Y_train_encod=to_categorical(Y_train_,num_class)
83                 model = tf.keras.models.Sequential()
84
85                 model.add(Dense(percpt,kernel_initializer='random_uniform',
input_shape=input_shape, activation=activations[activ]))
86                 model.add(Dense(num_class, kernel_initializer='random_uniform',
activation='softmax'))
87
88                 model.compile( optimizer='Adam',loss='categorical_crossentropy',
metrics=['accuracy'])
89                 model.fit(X_train_,Y_train_encod, epochs=100, batch_size=32,
verbose=0)
90                 temp_y=model.predict(X_test_)
91                 y_pred_=np.argmax(temp_y, axis=1)
92                 tmt=to_categorical(Y_test_,num_class)
93                 y_test=np.argmax(tmt, axis=1)
94
95                 cm = confusion_matrix(y_test, y_pred_,labels=[0.0,1.0])
96                 temp=(cm[0,0]+cm[1,1])/cm.sum()
97                 err_fold.append(1-temp)
98                 acc_fold.append(temp)
99
100                 error_mat[idx][activ]=(np.mean(err_fold)) # choose the perceptron
based on the min of error
101                 accuracy_mat[idx][activ]=(np.mean(acc_fold))
102
103             # find the perceptron with minimum error or loss
104             opt_perceptron=np.where(error_mat==error_mat.flat[np.argmin(error_mat
)])
105             optimal_perceptron=percept_list[int(opt_perceptron[0][0])]
106
107             print(f'optimal number of neurons:{optimal_perceptron}')

```

```

108
109 input_shape = (features,)
110 print(f'Feature shape: {input_shape}')
111
112 model = tf.keras.Sequential()
113 model.add(Dense(optimal_perceptron, kernel_initializer='random_uniform',
114                 , input_shape=input_shape, activation='relu'))
114 model.add(Dense(1, kernel_initializer='random_uniform', activation='
115                 sigmoid'))
115 model.compile(loss='binary_crossentropy', optimizer='Adam', metrics=['
116                 accuracy'])
116 model.fit(X_train, Y_train, epochs=100, batch_size=32, verbose=1)
117
118 # # Test the model after training
119 test_results = model.evaluate(X_test, Y_test, verbose=1)
120 print(f'Test results - Loss: {test_results[0]} - Accuracy: {
121         test_results[1]}')
121 p_pred = model.predict(X_test)
122 p_pred = p_pred.flatten()
123 # extract the predicted class labels
124 y_pred = np.where(p_pred > 0.5, 1, 0)
125 cm=confusion_matrix(Y_test, y_pred)
126 normalized_cm=cm / cm.astype(float).sum(axis=1)
127 error = (1- (cm[0,0]+cm[1,1])/cm.sum())
128
129 print(f'minimum probability of error:{error}')
130 print(f'confusion matrix:{cm}')
131 print(f'Normalized confusion matrix:{normalized_cm}')
132
133
134 min1, max1 = X_test[:, 0].min()-1, X_test[:, 0].max()+1
135 min2, max2 = X_test[:, 1].min()-1, X_test[:, 1].max()+1
136 # define the x and y scale
137 x1grid = np.arange(min1, max1, 0.1)
138 x2grid = np.arange(min2, max2, 0.1)
139
140 xx, yy = np.meshgrid(x1grid, x2grid)
141
142 r1, r2 = xx.flatten(), yy.flatten()
143 r1, r2 = r1.reshape((len(r1), 1)), r2.reshape((len(r2), 1))
144 # horizontal stack vectors to create x1,x2 input for the model
145 grid = np.hstack((r1,r2))
146 yhat = model.predict(grid)
147 contour_val=[xx,yy,yhat]
148 return error_mat, contour_val, y_pred
149
150 def SVM(sample_type):
151
152     grid = dict(kernel=kernel, C=C_list, gamma= gamma_list )
153     cross_val = KFold(n_splits=10)
154     grid_object = GridSearchCV(estimator=SVC(), param_grid=grid, cv=
155                               cross_val, scoring='accuracy')
155     grid_result = grid_object.fit(X_train, Y_train.reshape(-1))
156     print("Best: %f using %s" % (grid_result.best_score_, grid_result.
157                               best_params_))
158
159     svc = SVC(kernel = 'rbf', C = 100, gamma=0.01 )

```

```

160 svc.fit(X_train,Y_train)
161 y_pred=svc.predict(X_test)
162
163 cm=confusion_matrix(Y_test, y_pred)
164 normalized_cm=cm / cm.astype(float).sum(axis=1)
165 error = (1- (cm[0,0]+cm[1,1])/cm.sum())
166
167 print(classification_report(Y_test, y_pred))
168 print('Model accuracy score : {0:0.4f}'.format(accuracy_score(Y_test,
169 y_pred)))
170
171 grid_temp=list(grid_result.cv_results_['mean_test_score']*100)
172
173 min1, max1 = X_test[:, 0].min()-1, X_test[:, 0].max()+1
174 min2, max2 = X_test[:, 1].min()-1, X_test[:, 1].max()+1
175
176 x1grid = np.arange(min1, max1, 0.1)
177 x2grid = np.arange(min2, max2, 0.1)
178
179 xx, yy = np.meshgrid(x1grid, x2grid)
180
181 r1, r2 = xx.flatten(), yy.flatten()
182 r1, r2 = r1.reshape((len(r1), 1)), r2.reshape((len(r2), 1))
183
184 grid = np.hstack((r1,r2))
185 yhat = svc.predict(grid)
186 contour_val=[xx,yy,yhat]
187
188 return contour_val,y_pred,grid_temp
189
190 def plot_results(sample_type,classifier_type):
191
192
193 if classifier_type=='MLP':
194     error_matt,contour_val,y_pred=MLP(sample_type)
195
196     #plot min error vs number of perceptrons
197
198     tt=np.array(percept_list)
199     bar1=error_matt[:,0]
200
201     fig = plt.figure(figsize = (10, 7))
202     ax2 = plt.axes()
203     ax2.scatter(tt,bar1)
204     ax2.plot(tt,bar1)
205
206     plt.legend(activations)
207     plt.xlabel('Number of epochs')
208     plt.ylabel('Average minimum Probability of Error')
209     plt.title('Minimum Error for different number of epochs')
210     plt.show()
211
212 elif classifier_type=='SVM':
213     contour_val,y_pred,grid_temp=SVM(sample_type)
214
215
216     # accuracy percentage SVM

```

```

217
218
219     plt.plot(gamma_list,grid_temp[0:6])
220     plt.plot(gamma_list,grid_temp[6:12])
221     plt.plot(gamma_list,grid_temp[12:18])
222     plt.plot(gamma_list,grid_temp[18:24])
223     plt.plot(gamma_list,grid_temp[24:30])
224     plt.plot(gamma_list,grid_temp[30:36])
225     plt.xscale('log')
226
227     plt.xlabel("Values of Gamma")
228     plt.ylabel("Accuracy in %")
229     plt.legend(['C= 1000','C = 100','C = 10','C = 1.0','C = 0.1','C =
0.01'])
230     plt.title("Accuracy for all combinations of C and gamma")
231     plt.show()
232 # actual data distribution for test data
233 x0=[X_test[i] for i in range(samples[1]) if Y_test[i]==0]
234 x1=[X_test[i] for i in range(samples[1]) if Y_test[i]==1]
235
236 s1=[2 for i in range(len(x0))]
237
238 s2=[2 for i in range(len(x1))]
239
240
241 plt.scatter((np.array(x0))[:,0],(np.array(x0))[:,1],s1)
242 plt.scatter((np.array(x1))[:,0],(np.array(x1))[:,1],s2)
243
244 plt.legend(['class -1','class +1'])
245 plt.title('Actual Data Distribution')
246 plt.show
247
248 # contour plot of decision boundary and the classified data
249 x_temp=contour_val[0]
250 y_temp=contour_val[1]
251
252
253 x00t = [i for i in range(10000) if (Y_test[i] == 0 and y_pred[i] == 0)
254 ]
255 x01t = [i for i in range(10000) if (Y_test[i] == 0 and y_pred[i] == 1)
256 ]#fp
257 x10t = [i for i in range(10000) if (Y_test[i] == 1 and y_pred[i] == 0)
258 ]#fN
259 x11t = [i for i in range(10000) if (Y_test[i] == 1 and y_pred[i] == 1)
260 ]
261
262 plt.contourf(x_temp, y_temp, contour_val[2].reshape(x_temp.shape),
263             cmap='gist_heat')
264 plt.plot(X_test[x00t,0],X_test[x00t,1],'_',color='g', markersize =
265 1.95)#g
266 plt.plot(X_test[x01t,0],X_test[x01t,1],'_',color='r', markersize =
267 1.95)#r
268 plt.plot(X_test[x11t,0],X_test[x11t,1],'x',color='g', markersize =
269 1.85)#g
270 plt.plot(X_test[x10t,0],X_test[x10t,1],'x',color='r', markersize =
271 1.85)#r
272
273
274

```



```

265 plt.legend(['class 0 correctly labeled','class 0 wrongly labeled',
266             'class 1 correctly labeled','class 1 wrongly labeled'])
267 plt.xlabel("x1")
268 plt.ylabel("x2")
269 plt.show()
270 def main():
271     plot_results(samples[0], 'MLP')
272     plot_results(samples[0], 'SVM')
273
274 if __name__ == "__main__":
275     activations=['relu']
276     num_folds=10
277     num_class=2
278     features=2
279     samples=[1000,10000]
280     m1_n=np.zeros(2)
281     m2_n=np.zeros(2)
282
283     c1_n=np.eye(2)
284     c2_n=np.eye(2)
285
286     theta_margin=[-(np.pi),(np.pi)]
287     mu_vector=[m1_n,m2_n]
288     sigma_vector=[c1_n,c2_n]
289     prior=np.array([0.5,0.5])
290     percept_list=np.logspace(1,3,num = 15,endpoint = True,base = 5,dtype =
291                          int)
292     percept_list=percept_list.tolist()
293
294     kernel = ['rbf'] # can add multiple kernels
295     C_list = [1000, 100, 10, 1.0, 0.1, 0.01]
296     gamma_list = [0.01,0.1, 1,10,100,1000]
297     data_train,label_train=data(samples[0])
298     data_test,label_test=data(samples[1])
299
300     X_train=data_train.T
301     X_test=data_test.T
302
303     Y_train =label_train.T
304     Y_test = label_test.T
305
306     X_train = X_train.reshape(X_train.shape[0], features)
307     X_test = X_test.reshape(X_test.shape[0], features)
308
309     main()

```

Listing 1: Question 1

```

1 # Question 2
2 import numpy as np
3 import random
4 import matplotlib.pyplot as plt
5 from sklearn.mixture import GaussianMixture
6 from sklearn.model_selection import cross_val_score
7 from scipy . stats import multivariate_normal as mvn
8 import warnings
9 warnings.filterwarnings("ignore", category=RuntimeWarning)
10 import torchvision.transforms as transforms
11
12 from sklearn.preprocessing import minmax_scale
13 import cv2
14
15 import matplotlib.pyplot as plt
16 import matplotlib.cm as cm
17 import numpy as np
18
19
20 def GMM():
21
22     num_fold=10
23     gmm_list=[1,2,3,4,5,6,8,10]
24     data_image = cv2.imread('/content/drive/MyDrive/119082.jpg')
25     feature_vector=np.zeros((data_image.shape[0] * data_image.shape[1],5))
26     for i in range(data_image.shape[0]):
27         for j in range(data_image.shape[1]):
28             feature_vector[i*data_image.shape[1] + j,0]=i
29             feature_vector[i*data_image.shape[1] + j,1]=j
30             feature_vector[i*data_image.shape[1] + j,2]=data_image[i,j,2]
31             feature_vector[i*data_image.shape[1] + j,3]=data_image[i,j,1]
32             feature_vector[i*data_image.shape[1] + j,4]=data_image[i,j,0]
33
34     min_max=minmax_scale(feature_vector)
35     gmm_avg=np.zeros((len(gmm_list)))
36     gmm_var=np.zeros((len(gmm_list)))
37     temp=[]
38     for idx,i in enumerate(gmm_list):
39         gmm=GaussianMixture(i,covariance_type='full',random_state=None)
40         scor=cross_val_score(gmm,min_max,cv=num_fold) # this is the log
41         avg_scor=np.mean(scor)
42         var_scor=np.std(scor)
43
44         gmm_avg[idx]=avg_scor
45         gmm_var[idx]=var_scor
46
47         temp.append(avg_scor)
48
49     model_sel=np.argmax(temp)
50     model_sel=gmm_list[model_sel]
51     vv=np.array(temp)
52
53
54     plt.plot(gmm_list,vv,c='b', mfc='red',marker='o')
55     plt.xlabel('model order')
56     plt.ylabel(f'log likelihood')
57     plt.title(f'log likelihood under different GMM model orders',fontsize

```

```

    =12.)
58 plt.show()
59
60 gmm_model=GaussianMixture(model_sel,covariance_type='full').fit(
    data_image.reshape((-1,3)))
61 gmm_labels=gmm_model.predict(data_image.reshape((-1,3)))
62
63 segmented=gmm_labels.reshape((data_image.shape)[0],(data_image.shape)
    [1])
64 cv2.imwrite('segment.jpg',segmented)
65 plt.imshow(np.array(segmented))
66
67 GMM()

```

Listing 2: Question 2